

PowerModels Strategic Options Analysis

Generated: 2026-03-08 Sources: Planning vault (23 facets, 11 audiences, 29 intersections, 8 source documents), Implementation vault (scan-data.json, 10 architecture docs, 6 artifact docs), vault-project-analysis-2026-03-08-v3.md, mitigation-plan-2026-03-08.md, accelerated-mitigation-plan-2026-03-08.md, business-plan.md, strategic_summary_current.md, investor-market-evaluation.md

Context

PowerModels faces a strategic decision point. The company has:

- **A validated thesis:** Tax-prep-driven entry, firm-first GTM, 85/15 boundary, event-sourced architecture
- **A working product:** 174 stories shipped, 5 milestones, AI pipeline, local-first WPF+Excel dual interface
- **Early market validation:** 3 design partner firms (Liberty Tax, Stewart, Bancroft) have validated the need; all require a working demo/showcase before moving to initial trials
- **Technical debt:** God aggregate (3,557 lines, 87 Apply methods), 0/27 RestoreFromEvents tests, 46 RMs replaying on startup, no process managers, no checkpoint persistence
- **A cloud engineer** (Joao) working on cloud/security infrastructure independently of the mitigation plan
- **A detailed mitigation plan** for the existing codebase, compressible to 5–6 weeks with expert team + agents
- **Team on sweat equity** for ~2 years, collectively owning 12% of the company

Three courses of action are under consideration. This analysis evaluates each against what investors actually underwrite and what the business needs to succeed.

The Investor Lens

Before evaluating options, we must establish what investors at each stage actually look for (sourced directly from the planning vault audience profiles and investor-market-evaluation.md):

Angel (\$500K) — What They Underwrite

- **Founder quality and commitment** — team working for sweat equity signals conviction
- **Problem-solution fit** — is this a real, painful problem worth solving?
- **Early traction** — customer interest, pilots, LOIs
- **Capital discipline** — scrappiness, small checks, long runway
- **Unique insight** — what do they see that others don't?
- **Path to next milestone** — can they hit proof points for seed?

Seed/Series A (\$3–5M) — What They Underwrite

- **Product-market fit evidence** — 3–5+ reference customers, ideally with expansion
- **Go-to-market repeatability** — can the company acquire customers predictably?
- **Unit economics** — CAC/LTV, payback period, does it improve with scale?
- **Competitive moats** — what's defensible?
- **Team completeness and execution** — can this team scale?
- **Capital efficiency** — how much capital to reach next inflection point?
- **Path to Series B** — line-of-sight to \$5–10M ARR

What Gets Punished (2026 Market)

- "AI will do accounting" positioning (credibility gap)
- Services creep without software margins

- Heroic onboarding that won't scale
- Past dates without status updates (instant doubt)
- Valuation disconnected from evidence
- **Thin SaaS wrappers** — investors increasingly doubt any SaaS whose value can be replicated by foundation models (Anthropic, OpenAI)

The "Anthropic Guts SaaS" Question

There is growing market doubt about whether SaaS offerings survive foundation model disruption. Investors are asking: *"What stops Claude/GPT from replacing this product?"*

SaaS products that are vulnerable:

- Thin wrappers around LLM calls (summarization, content generation, chatbot builders)
- Workflow tools whose logic is prompt-engineerable
- Data transformation tools with no proprietary data or domain model
- Products where the AI IS the product

What makes PowerModels structurally resistant:

Defense Layer	Why It Can't Be Replicated by a Foundation Model
Event-sourced financial record	An append-only, time-ordered ledger with deterministic replay is a data architecture, not an AI task. No prompt produces this.
27 aggregates encoding accounting domain invariants	Business rules enforced at the aggregate level (balanced entries, period boundaries, fiscal year constraints) require a persistent, stateful system — not a stateless inference call.
85/15 boundary	The explicit design choice to stop at 85% and preserve professional judgment is the opposite of "AI will do everything." It's a trust architecture.
Reconciliation elimination	Single source of truth across financial records is a structural property of the data model, not a feature that can be prompted.
Regulatory accountability	Tax preparation requires audit trails, deterministic outputs, and professional sign-off. LLMs are probabilistic — they cannot serve as the system of record.
Firm-first distribution	The GTM is institutional adoption, not consumer self-serve. Firms don't adopt LLM prompts as infrastructure.

PowerModels uses AI correctly — as a tool inside a trusted architecture:

- AI classifies transactions (the 85% mechanical work)
- AI extracts data from shoebox documents
- AI suggests mappings in the reconciliation pipeline
- But the financial record, the accounting rules, the audit trail, and the professional review workflow are all deterministic, event-sourced, and human-governed

The investor pitch on this point: *"We use AI to accelerate the mechanical 85% of accounting preparation. But the financial record itself — the thing firms trust, regulators audit, and lenders rely on — is a deterministic, event-sourced system that no foundation model can replace. We're not a wrapper around AI. We're an accounting infrastructure platform that uses AI as one input among many."*

This positioning is strengthened, not threatened, by foundation model advances. Better models make the 85% faster and cheaper while the 15% professional judgment layer and the underlying architecture remain PowerModels' moat.

Option A: Fix What We Have

Description

Execute the accelerated mitigation plan on the existing codebase. 5–6 weeks with the expert team (2 ES experts, principal, 3 support) + AI agents. Cloud migration continues in parallel (Joao). Ship cloud-based product for April 2026 GTM.

What This Preserves

Asset	Status	Investor Value
174 shipped stories across 5 milestones	Proven	Execution evidence
Event-sourced architecture (ReactiveDomain)	Working	Technical moat — incumbents can't retrofit
WPF + Excel dual interface	Working	Professional-grade UI for tax accountants
AI pipeline (Semantic Kernel, gpt-4o)	Working	AI empowerment positioning
Shoebox-to-trial-balance pipeline	Validated (tested against disorganized client records)	Core value prop demonstration
Bank statement PDF ingestion	Working	Completeness proof artifact
Reconciliation pipeline (MappingStep → RuleStep → AIStep → FormatConversion)	Working (with defects)	End-to-end preparation cycle
Design partner pipeline (3 firms)	Active	Angel round proof
QB/Xero import capability	Working	Onboarding scalability answer
85/15 boundary implementation	Working	"Unusually smart" investor advantage
Domain knowledge embedded in 27 aggregates, 80 read models	5+ years of work	Irreplaceable institutional knowledge

What the Mitigation Fixes

Problem	Fix	Timeline
No RestoreFromEvents tests → replay safety risk	30/30 coverage	Weeks 1–5

God aggregate (3,557 lines) → compound defect risk	SFM extraction: 87 → ~67 Apply, 31 → ~11 events	Weeks 4–5
No checkpoint persistence → startup replay amplification	46/46 RMs checkpointed, <2s startup at 15K events	Weeks 1–5
No process managers → implicit sagas → 29% reconciliation defect rate	2 explicit PMs (Reconciliation, PeriodClose), <10% defect rate	Weeks 3–5
Missing accounting aggregates (Journal, ReconciliationState, PeriodClose)	3 new aggregates designed and implemented	Weeks 2–5
No caching for SFM GetById	CachingRepository with persisted snapshot tier	Week 2

Investor Scorecard

Investor Criterion	Score	Evidence
Product-market fit evidence	★★★★☆	3 firms (Liberty, Stewart, Bancroft) validated need; all awaiting working demo for next step; shoebox pipeline proven
Go-to-market repeatability	★★★★☆	Firm-first model defined; 3 firms validated need but all gated on demo; association channel not yet proven at scale
Unit economics	★★★★☆	\$20/mo validated at micro tier; \$48K–\$87K ARR/firm modeled; CAC per firm not per client
Competitive moats	★★★★★	Event-sourced architecture, 85/15 boundary, reconciliation elimination — all validated by code
Team execution	★★★★★	5+ years together, ES/DDD/CQRS experts, 174 stories shipped, detailed mitigation plan shows self-awareness
Capital efficiency	★★★★★	~2 years on sweat equity, detailed plan shows exactly what capital buys
Path to next round	★★★★☆	Clear: mitigate → cloud launch → design partners → seed metrics
Time to revenue	★★★★☆	April 2026 cloud launch → tax season → first revenue Q2–Q3 2026
Technical risk	★★★★☆	God aggregate and implicit sagas are real risks; mitigation plan addresses them directly

Risks

- **Mitigation plan execution risk:** Plan is detailed but still 5–6 weeks of intensive work; any delay pushes cloud launch
- **Cloud migration coordination:** Must align mitigation output with James's cloud infrastructure work
- **Technical debt may have surprises:** SFM RestoreFromEvents test (87 Apply methods) may reveal >5 bugs requiring event versioning
- **Sweat equity fatigue:** Team has been on deferred compensation for ~2 years; another 5–6 weeks of intense execution before any revenue signal

Capital Required

- **Angel round:** ~\$500K (already planned)
- **Use:** Team continuity through mitigation + cloud launch + design partner onboarding
- **Runway to:** Seed metrics (10+ firms, expansion data, CAC payback evidence)

Option B: Rewrite a Smaller Product From Scratch

Description

Build a new, focused product from scratch — likely targeting the shoebox-to-trial-balance pipeline only, or just the micro/1099 tier. Cloud-native from day one. Potentially use a modern stack (web-based, not WPF).

What This Abandons

Asset Lost	Replacement Cost	Impact
174 shipped stories (5 milestones)	12–24 months of rewrite	All execution evidence gone
Event-sourced architecture (ReactiveDomain)	Must re-implement or choose different architecture	Technical moat weakened or abandoned
27 aggregates modeling accounting domain	6–12 months domain modeling	5+ years of domain knowledge lost
80 read models (46 ReadModelBase + 34 TransientSubscriber)	3–6 months to rebuild	Query layer expertise lost
AI pipeline (Semantic Kernel integration)	2–4 months rebuild	Delay to AI empowerment story
Reconciliation pipeline (4-step, working)	2–3 months rebuild	Core value prop delayed
WPF + Excel dual interface	Decision: rebuild in WPF? or pivot to web?	Either way, months of UI work
Design partner validation (3 firms)	Must re-validate with new product	Pipeline reset
QB/Xero import pipeline	Must rebuild	Onboarding story delayed
ProcessManager infrastructure (ready for integration)	Must re-implement	Infrastructure cost repeated
Checkpoint/persistence infrastructure (ready)	Must re-implement	Same

What This Gains

Benefit	Reality Check
Clean architecture from day one	But the existing architecture IS clean (ES/DDD/CQRS); the debt is in one god aggregate and missing tests — fixable in 5–6 weeks

Cloud-native from day one	Cloud migration is already underway with Joao; the domain layer ports cleanly
Smaller scope = faster MVP	But MVP already exists — 3 firms have validated the need and are waiting for a working demo
Web-based = broader accessibility	True, but WPF is correct for tax accountant workflow; web is Phase 2+
Modern stack appeal	The current stack (C#, ReactiveDomain, EventStoreDB, Semantic Kernel) IS modern and purpose-built

The Fundamental Problem With Option B

You already have the MVP. Three firms have validated the need and are waiting for a working demo/showcase to move to initial trials. The shoebox pipeline works. The reconciliation pipeline works (with fixable defects). The 85/15 boundary is implemented.

A rewrite doesn't produce a better product faster — it restarts the clock on:

- Domain modeling (the hardest part — 5+ years of accounting domain knowledge)
- Design partner trust (3 firms waiting for demo — a rewrite resets that clock)
- Proof artifacts (all investor evidence resets)
- Team momentum (morale hit from discarding 2 years of sweat equity work)

Investor Scorecard

Investor Criterion	Score	Evidence
Product-market fit evidence	★★☆☆☆	Resets to zero — prior validation doesn't transfer
Go-to-market repeatability	★★☆☆☆	No product to sell; GTM timeline pushed 6–12+ months
Unit economics	★★☆☆☆	Theoretical only — no real data possible for months
Competitive moats	★★★☆☆	Team expertise preserved, but architectural moat must be rebuilt
Team execution	★★★☆☆	Demonstrates willingness to make hard calls, but also signals prior execution failure
Capital efficiency	★★☆☆☆	~2 years of sweat equity effectively written off; capital now funds rebuild not growth
Path to next round	★★☆☆☆	Pushed 9–15 months — must rebuild → re-validate → re-prove
Time to revenue	★★☆☆☆	12–18 months minimum before design partner readiness
Technical risk	★★★★☆	Cleaner start, but new risks: will the team ship faster the second time? History says no

The Rewrite Trap

The decision to rewrite is often driven by the feeling that "if we just started over, we'd do it right." But the evidence from the implementation vault shows:

1. **The architecture IS right.** ES/DDD/CQRS with ReactiveDomain is exactly the correct architecture for this domain. The problems (god aggregate, missing tests, no checkpoints) are fixable without rewriting.
2. **The domain model IS right.** 27 aggregates, 80 read models, domain events correctly modeled — this represents 5+ years of accounting domain expertise encoded in code.
3. **The problems are specific and addressable.** The v3 analysis identified exactly what's wrong and the mitigation plan shows exactly how to fix it in 5–6 weeks.

A rewrite says "we'll throw away what works to avoid fixing what's broken" — when fixing what's broken is faster than rewriting what works.

Capital Required

- **Angel round:** ~\$500K–\$750K (more needed because timeline extends)
- **Additional funding before revenue:** Likely another \$500K–\$1M to reach design partner stage
- **Total to seed-ready:** \$1–1.5M vs \$500K for Option A
- **Time to seed-ready:** 12–18 months vs 4–6 months for Option A

Option C: Slim the Team + Seek Market Validation

Description

Reduce the team to a skeleton crew. Pause or drastically slow product development. Focus on market validation — customer interviews, firm conversations, possibly exploring Canadian market with a new team. Use minimal capital to extend runway while seeking stronger product-market fit signals.

What This Preserves

Asset	Status	Investor Value
Existing codebase (frozen)	Available for future use	Retains option value
Business plan + strategic frameworks	Intact	Can be pitched to new investors
Design partner relationships	At risk — all 3 firms (Liberty, Stewart, Bancroft) need a working demo to proceed; pausing loses them	Fragile
Founding team (Chris, Lisa)	Preserved	Core IP holders retained
Market research (12 interviews)	Available	Validation artifacts available

What This Risks

Risk	Severity	Explanation
Design partner pipeline collapse	CRITICAL	All 3 firms (Liberty, Stewart, Bancroft) have validated the need but require a working demo/showcase to move to initial trials. Pausing development loses the April 15 tax deadline window entirely.
Team dissolution	HIGH	Development team (12% equity, ~2 years sweat) will leave for paid work. ES/DDD/CQRS expertise is extremely rare — this team cannot be

		reassembled.
Competitive window closes	HIGH	Industry is converging on AI empowerment (not replacement) positioning per Accounting Podcast. Other entrants will fill the gap.
April 2026 cloud launch missed	CERTAIN	No development = no cloud launch = no GTM = no revenue signal
Sweat equity devalued	HIGH	12% ownership for ~2 years of work becomes worthless if company stalls
Canadian restart fallacy	HIGH	Rebuilding with a new team means losing the ES/DDD/CQRS expertise that IS the moat

The "Further Market Validation" Question

The planning vault already contains strong validation signals:

What's validated:

- Tax preparation as forcing function (Liberty, Stewart, Bancroft all confirm)
- 85/15 boundary resonates with practitioners (Accounting Podcast industry consensus)
- Shoebox pain is acute and universal (Stewart: "100 hours unbillable setup")
- Firm-first model makes economic sense (CAC per firm, not per client)
- \$20/mo pricing validated at micro tier
- All 3 firms need a working demo/showcase to proceed to trials — signals demand contingent on product readiness, not market doubt
- Event-sourced architecture addresses the core structural problem

What's not yet validated:

- Full tax season throughput with the product
- Expansion mechanics (clients per firm over time)
- Onboarding scalability (time-to-first-cycle compression)
- Pricing at Solo/Partnership and SMB tiers
- Association channel conversion rates

The gap is **execution validation, not market validation**. The market is saying yes. The question is whether the product can deliver at the quality and speed firms require. Slimming the team makes that harder, not easier.

Investor Scorecard

Investor Criterion	Score	Evidence
Product-market fit evidence	★★☆☆☆	Existing interviews preserved, but no new proof generated; active pipeline at risk
Go-to-market repeatability	★☆☆☆☆	No product in market; GTM frozen
Unit economics	★☆☆☆☆	Theoretical only; no path to real data without product
Competitive moats	★★☆☆☆	Moat exists in the team's heads but is not being built into defensible IP
Team execution	★☆☆☆☆	Signals failure to execute with current team; "we need to find a different team" is a red flag

Capital efficiency	★★★★☆	Low burn, but no progress per dollar — capital buys time, not value
Path to next round	★☆☆☆☆	No clear milestones; "seeking further validation" is not a fundable plan
Time to revenue	★☆☆☆☆	Indefinite
Technical risk	★★★★☆	Code freezes don't improve; they rot. Returning to a frozen codebase is harder than maintaining a live one

Capital Required

- **Skeleton crew:** \$100–200K (low burn, founder salaries only)
- **Time to decision:** 3–6 months of conversations
- **Total to restart:** If Canadian team, \$500K–\$1M to rebuild + \$500K to operate = \$1–1.5M
- **Time to market after restart:** 18–24 months (new team, new codebase, re-validate)

Side-by-Side Comparison

Timeline to Investor Milestones

Milestone	Option A (Fix)	Option B (Rewrite)	Option C (Slim)
Product in market	6–8 weeks (mitigation + cloud launch)	9–15 months	Indefinite
Design partner validation	Q2 2026 (3 firms awaiting demo)	Q4 2026–Q1 2027	At risk of loss
Angel round proof points	April 2026	Q1–Q2 2027	None generated
Seed round readiness	Q3–Q4 2026	Q2–Q4 2027	2027+ (if restart)
First revenue	Q2–Q3 2026	Q4 2026–Q1 2027	Indefinite
Path to \$3–4M ARR	12–18 months post-launch	24–36 months	Unknown

Capital Efficiency

Metric	Option A	Option B	Option C
Capital to angel proof points	\$500K	\$750K–\$1M	\$100–200K (but no proof points)
Capital to seed readiness	\$500K	\$1–1.5M	\$1–1.5M (restart)
Sweat equity preserved	✅ Full value	❌ Partially destroyed	⚠️ At risk of total loss
Prior work leveraged	✅ 174 stories, 5+ years domain	❌ Started over	❌ Frozen/abandoned

Team retention	✔ Motivated by clear plan	⚠ Morale risk from discard	✖ Team dissolves
----------------	---------------------------	----------------------------	------------------

Investor Narrative Strength

Element	Option A	Option B	Option C
"We shipped"	✔ 174 stories, design partner validated	✖ "We're rebuilding"	✖ "We're pausing"
"Customers want it"	✔ 3 firms validated need, awaiting demo to start trials	⚠ "They wanted the old version"	⚠ "We think they still do"
"We know the problem"	✔ Self-aware analysis + fix plan	⚠ "We had to start over"	✖ "We're not sure yet"
"We're capital efficient"	✔ 2 years sweat → working product	✖ 2 years sweat → discarded	✖ 2 years sweat → frozen
"Here's the moat"	✔ ES architecture in production	⚠ "We'll rebuild the moat"	✖ "The moat is in our heads"
"Here's the timeline"	✔ 5–6 weeks to fix → April launch	⚠ 12–18 months to MVP	✖ "We need time to figure it out"

The Investor Narrative for Each Option

Option A: "We Built It, We Know What's Wrong, We're Fixing It"

"We have a validated product with a design partner on hold for our stable demo, and a detailed, evidence-based plan to fix specific technical debt in 5–6 weeks. Our expert team — 5+ years together, world-class in the exact architecture we use — will execute this while cloud deployment continues in parallel. We'll be in market by April 15, 2026 — the tax filing deadline — with 3 design partner firms in pipeline."

Investor reaction: This is fundable. Self-aware founders who built something real, identified problems honestly, and have a credible fix plan. The team's domain expertise (ES at Event Store, financial systems at Wellington/Linedata/Hanover) directly addresses the technical risk. Capital buys execution, not speculation.

Option B: "We Built It Wrong, We're Starting Over"

"We spent 2 years building a product, but we've decided to start over with a smaller scope. We believe we can build it better the second time."

Investor reaction: Red flag. Why did it take 2 years to realize? What ensures the rewrite succeeds? The team already knows ES/DDD/CQRS — if the current architecture is wrong, what changes? If the architecture is right (it is), why rewrite? This narrative destroys the execution credibility that's one of the company's strongest assets.

Option C: "We Need More Time to Validate"

"We're slimming down to seek further market validation, possibly with a new Canadian team."

Investor reaction: Fatal. After 2 years of development and active beta users, asking for more validation time signals either: (a) the market isn't there (contradicts the evidence), (b) the team can't execute (contradicts the track record),

or (c) the founders have lost conviction. None of these are fundable.

The Cloud Engineer Factor

Joao is working on cloud hosting/security infrastructure independently. This is a **critical asset for Option A**:

- Cloud migration scope: Event store → cloud-hosted persistence, multi-tenant (firm = tenant), auth, cloud-to-cloud QB/Xero integration
- The domain layer (aggregates, events, read models, services) ports to cloud without change — ES/DDD/CQRS is deployment-agnostic
- The mitigation plan output (fixed aggregates, checkpointed RMs, explicit PMs) makes the cloud deployment more reliable, not harder
- **Two parallel tracks converge**: Mitigation team fixes the domain → Joao makes it cloud-deployable → April launch

For Options B and C, Joao's cloud work is either:

- **Option B**: Partially wasted (new codebase may need different cloud architecture)
 - **Option C**: Completely wasted (nothing to deploy)
-

The Canadian Team Question

Option C suggests possibly "restarting with a brand new Canadian team." This deserves direct analysis:

What the current team brings that cannot be replaced

1. **ES/DDD/CQRS expertise at production scale**: Chris (CTO EventStore, Chief Architect at Linedata, Principal at PerkinElmer, Lead at Wellington), plus a development team who are "ES/DDD/CQRS experts" with 5+ years together. This concentration of event-sourcing expertise in an accounting product company is globally rare.
2. **5+ years of accounting domain modeling**: 27 aggregates, 80 read models, reconciliation pipeline, AI classification pipeline, shoebox ingestion — this domain knowledge is embedded in code and in the team's heads. A new team starts at zero.
3. **Existing customer relationships**: Liberty Tax, Stewart, and Bancroft know and trust this team. All three have validated the need and are waiting for a demo. Starting over with a new team resets these relationships.
4. **Sweat equity investment**: \$0 in cash compensation for ~2 years. This signals extraordinary conviction and commitment. A new team will require cash from day one.

What a Canadian team offers

- Canadian R&D tax credits (SR&ED) — significant (up to 35% of qualifying expenditures)
- Access to Canadian VC ecosystem (different investor pool)
- Lisa's Canadian network and fundraising experience (\$100M+ raised in Canada)
- Local market entry for Canadian accounting firms (~27,000 firms, CAD \$18–28B market)

The hybrid answer

The Canadian market advantages don't require abandoning the current team. They can be captured through:

- Incorporating in Canada while retaining the current development team
- Using SR&ED credits on the existing codebase development

- Targeting Canadian associations (equivalent of PSTAP) for GTM
 - Lisa's Canadian network for capital strategy
-

Recommendation

Option A is the clear choice, with elements of the other options captured within it.

Why:

1. **The fastest path to investor proof points.** April 2026 cloud launch with active design partners = angel round proof within weeks, not months or years.
2. **The most capital-efficient path.** \$500K buys a complete fix + cloud launch + design partner validation. Options B and C cost 2–3× as much to reach the same milestones.
3. **The only path that preserves existing assets.** 174 stories, 5+ years of domain knowledge, 3 design partner relationships, working AI pipeline, working shoebox pipeline — all preserved.
4. **The only path that retains the team.** The development team is on sweat equity. They need to see forward momentum and a path to compensation. Option A provides that. Options B and C risk losing them permanently.
5. **The self-awareness is itself an investor asset.** The v3 analysis, mitigation plan, and accelerated plan demonstrate extraordinary technical self-awareness. Investors fund teams that know their weaknesses and have credible plans to address them — not teams that throw away work or stall for "more validation."
6. **The technical debt is fixable, not fatal.** The problems are specific: one god aggregate, missing tests, no checkpoints, implicit sagas. The solutions are known: extraction, RestoreFromEvents, ICheckpointStore, ProcessManagerBase. The team that built these solutions is available. 5–6 weeks.

Capture Option C's market validation within Option A

- Run the accelerated mitigation (5–6 weeks) **while Lisa continues firm conversations**
- Use the cloud launch as a forcing function for Stewart and Liberty Tax conversion
- Use April 2026 as the angel round proof deadline
- Evaluate Canadian incorporation / SR&ED as a capital strategy decision independent of product execution

The one-sentence pitch

"We have a validated product with a design partner waiting for our stable demo, a world-class ES team that built it together over 5 years, a detailed 6-week plan to fix specific technical debt, and a cloud launch targeting April 15 — tax filing deadline — with 3 design partner firms ready to validate the tax-prep-to-85% workflow."

This is fundable on first read.

Decision Matrix Summary

Factor	Weight	Option A: Fix	Option B: Rewrite	Option C: Slim
Time to market	HIGH	★★★★★ (6–8 weeks)	★★★☆☆ (12–18 months)	★★☆☆☆ (indefinite)

Capital efficiency	HIGH	★★★★★ (\$500K)	★★★★☆ (\$1–1.5M)	★★★★☆ (low burn, no value)
Asset preservation	HIGH	★★★★★ (everything)	★★★★☆ (team only)	★★★★☆ (frozen code)
Team retention	HIGH	★★★★★ (clear plan)	★★★★☆ (morale risk)	★★★★☆ (dissolution)
Investor narrative	HIGH	★★★★★ (self-aware builders)	★★★★☆ (restart signal)	★★★★☆ (stall signal)
Design partner pipeline	HIGH	★★★★★ (3 firms active)	★★★★☆ (reset)	★★★★☆ (at risk)
Technical risk mitigation	MED	★★★★☆ (specific plan)	★★★★☆ (clean start)	★★★★☆ (no progress)
Market timing	HIGH	★★★★★ (April 2026)	★★★★☆ (2027)	★★★★☆ (window may close)
Canadian market access	LOW	★★★★☆ (via incorporation)	★★★★☆ (via new team)	★★★★★ (native)
TOTAL (weighted)		★★★★★	★★★★☆	★★★★☆

Appendix: What "Fix What We Have" Actually Means

For precision, the mitigation plan addresses these specific items (not "general cleanup"):

Item	Current State	Post-Mitigation State	Why It Matters for Investors
RestoreFromEvents tests	0/27	30/30	Proves ES architecture works correctly — not just claimed
SFM god aggregate	3,557 lines, 87 Apply	~1,500–2,000 lines, ~67 Apply	Removes #1 defect source; shows engineering discipline
RM startup replay	Full replay from position 0	Checkpoint-based resume (46/46)	Startup <2s at 15K events — deployable at scale
Process managers	0 (7 implicit sagas)	2 explicit (Recon, PeriodClose)	Reconciliation defect rate from 29% → <10%
Accounting domain	Missing Journal, ReconciliationState, PeriodClose	3 new aggregates designed and implemented	Foundation for accountant workflow features

FinancialModelService	166 commands in one handler	Split: ~100 core + ~60 mapping	Reduces coordination complexity
Caching	No aggregate snapshots	CachingRepository for SFM	GetById: full replay → <100 events from snapshot

Total cost: **5–6 weeks** with the team described. Total capital: **\$500K angel round** (already planned).

None of this requires changing the architecture, the domain model, the product scope, or the team. It requires executing a plan that the team themselves wrote, using infrastructure they themselves designed.